

# What is Neural Network Verification?

Yulia Alexandr

Foundations and Trends® in  
Programming Languages  
7:1-2

## Introduction to Neural Network Verification

Aws Albarghouthi

**now**  
the essence of knowledge

# Introduction to Neural Network Verification

by Aws Albarghouthi

A foundational text covering:

- Formal definitions of verification
- Methods and tools
- Applications to safety-critical systems

Available on [arXiv](#).

# Motivation

- Neural networks power critical applications:
  - ▶ Autonomous driving
  - ▶ Medical diagnosis
  - ▶ Defense and aerospace
- But they are black boxes:
  - ▶ Why did the model make a decision?
  - ▶ Can we guarantee it will not fail?

## Definition and Scope

Neural network verification is the problem of proving that a neural network satisfies certain desirable properties, such as:

- **Correctness:** Does the output match the intended specification?
- **Safety:** Does it avoid unsafe behaviors under all valid inputs?
- **Robustness:** Do small input perturbations leave outputs unchanged?

# Example Property

## Image classifier for road signs

- **Input:** Image of a stop sign
- **Property:** Even with noise (graffiti, lighting, stickers), the output should still be “Stop”



# Key Trade-offs in Neural Network Verification

Neural network verification systems need to balance three aspects:

- **Soundness:** Can we trust that the answer is correct?
- **Completeness:** Can an answer be provided in many cases?
- **Scalability:** Can the solver handle large networks efficiently?

# Approaches to Neural Network Verification

- Constraint-based techniques (complete verification)
  - ▶ Represent the neural network and desired properties as a system of constraints
  - ▶ Solve the constraints to prove or disprove that the network satisfies the properties
- Abstraction-based techniques (approximate verification)
  - ▶ Instead of executing the network on a single input, operate on an infinite set of inputs
  - ▶ Show that all inputs in this set satisfy the desired properties
  - ▶ More scalable but may produce conservative (approximate) guarantees

# Complete Verification

## Encoding Correctness Properties

**Idea:** Reduce the verification problem to checking satisfiability of a logical formula.

**Formal structure:**

$$\{\text{precondition}\}, \quad r \leftarrow f(x), \quad \{\text{postcondition}\}$$

**Example: Binary image classifier**

- **Network:**  $f_G : \mathbb{R}^n \rightarrow \mathbb{R}^2$   
**Input:** grayscale image  $x \in [0, 1]^n$
- **Goal:** small perturbations do not change prediction  
**Precondition:**  $|x - c| \leq 0.1$  ( $c$  = original cat image)  
**Postcondition:**  $r_1 > r_2$  (probability of cat > probability of dog)

# DPLL Algorithm

**Question:** How do we check whether a logical formula is satisfiable?

DPLL (Davis–Putnam–Logemann–Loveland) algorithm:

- Developed decades ago for checking satisfiability of Boolean formulas
- Can be extended to handle first-order formulas over theories
- Forms the basis of modern SAT and SMT solvers:
  - ▶ SAT solvers: check satisfiability of propositional logic formulas
  - ▶ SMT solvers: check satisfiability modulo theories (e.g., numbers, arrays, bit-vectors)

# Simplex Algorithm

The **simplex algorithm** is used for solving conjunctions of linear inequalities (literals) over real numbers.

Steps in neural network verification:

- Initially handles linear constraints directly using the simplex method
- ReLUs introduce piecewise-linear behavior
- ReLUs are encoded as disjunctions and solved using a SAT solver
- Instead, we can extend the solver to handle ReLUs **natively**:
  - ▶ directly reason about ReLU activations as part of its computations, instead of converting them into another form that the solver can understand (like disjunctions for a SAT solver)

This combination allows verification of neural networks with piecewise-linear activations more effectively than treating ReLUs purely as Boolean constraints.

# Approximate verification

## Motivation:

- Complete verification encodes the network and correctness property as a formula in first-order logic
- Leads to NP-complete problems (e.g., satisfiability modulo linear real arithmetic / mixed-integer linear programming)
- Scaling to large neural networks is challenging

## Approximate verification:

- Uses overapproximation (abstraction) of the network semantics
- Can prove correctness properties, but failure does not imply violation
- Based on **abstract interpretation** (Cousot & Cousot, 1977)
- Provides a scalable, pragmatic way to verify large neural networks